

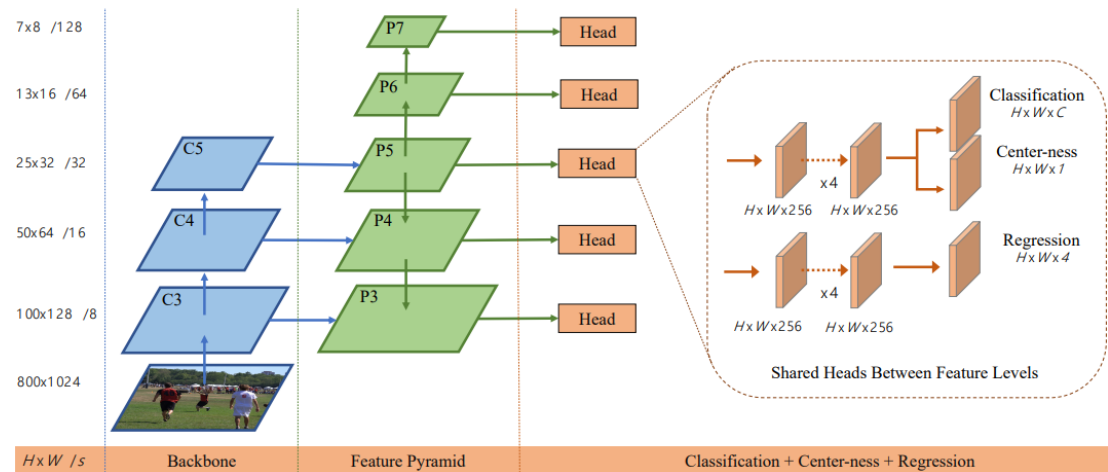

# EE898 Programming Assignment 1

FCOS: Fully Convolutional One-Stage Object Detection

Youngtaek Oh (youngtaek.oh [at] kaist.ac.kr)

# Programming Assignment 1

- In PA1, you will implement FCOS object detector, a fully convolutional anchor-free one-stage object detector.
- Please read the paper thoroughly to complete a baseline code and additional questions prepared in PA1.





## Detectron2



# Detectron2

<https://github.com/facebookresearch/detectron2>

## Dependencies

- We provide detectron2-based skeleton code, where detectron2 provides many convenient tools such as data processing, train/test loop, dataset evaluation, and visualization of results.
- We hope all of you to pay attention solely on writing codes about detection pipelines and FCOS head itself.
- So, in order to train your implemented FCOS detector, you need to install detectron2 on your environment.
- We recommend installing detectron2 on Linux OS, Python $\geq$ 3.6, Pytorch $\geq$ 1.3.

# Detectron2 Installation

- Requirements

- Linux or macOS with Python  $\geq 3.6$
- PyTorch  $\geq 1.3$
- torchvision that matches the PyTorch installation.
- OpenCV, optional, needed by demo and visualization
- pycocotools: `$ pip install cython; pip install -U 'git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI'`

- Build Detectron2 from source (recommended)

```
# install it from a local clone:  
$ git clone https://github.com/facebookresearch/detectron2.git  
$ cd detectron2 && python -m pip install -e .
```

<https://github.com/facebookresearch/detectron2/blob/master/INSTALL.md>

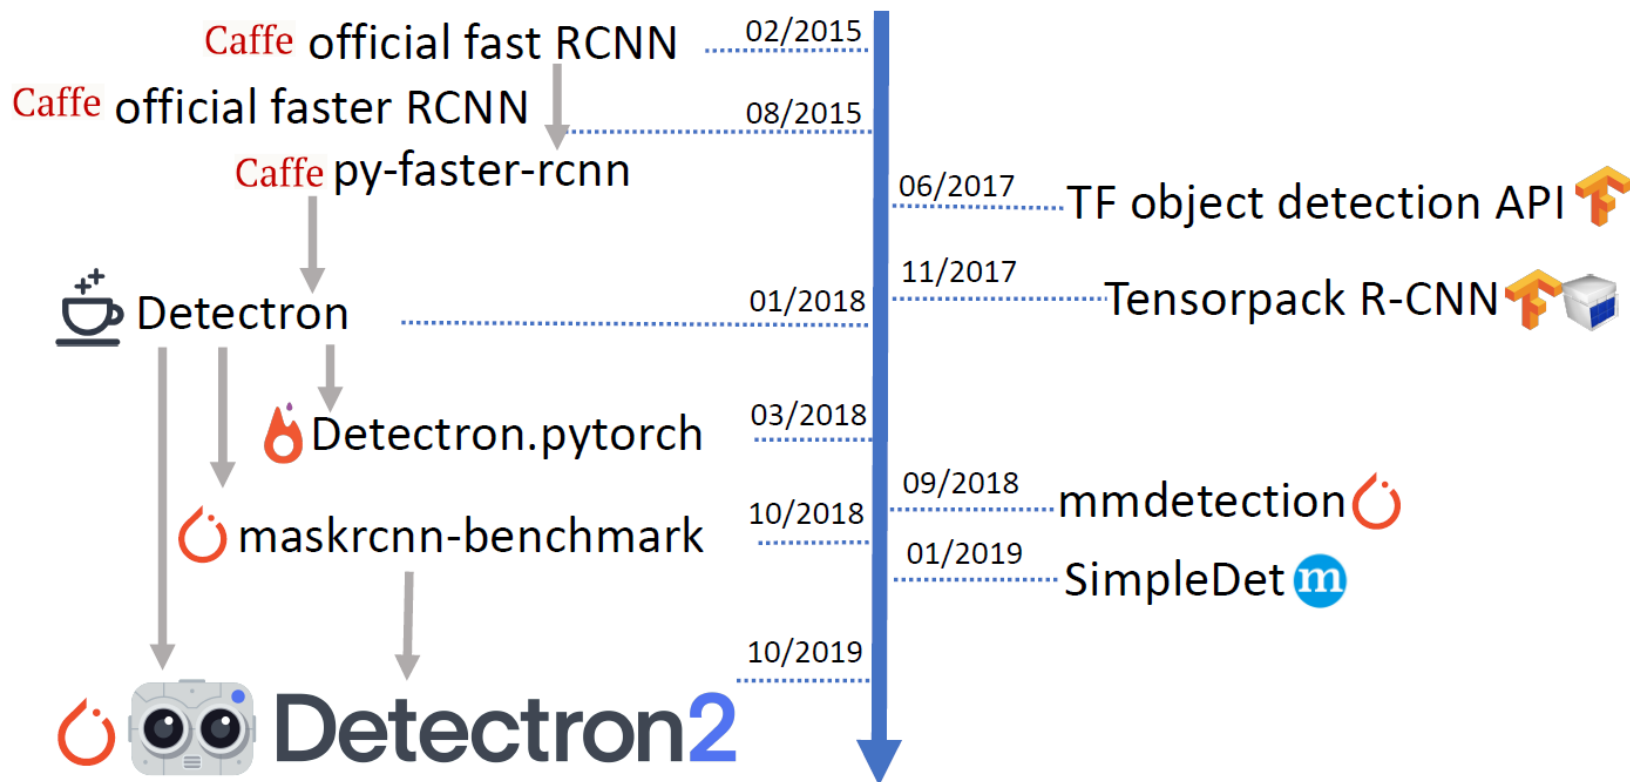
# Common installation issues

- Common Installation Issues
  - Undefined torch/aten/caffe2 symbols, or segmentation fault immediately when running the library.
  - Undefined C++ symbols (e.g. `GLIBCXX`) or C++ symbols not found.
  - "Not compiled with GPU support" or "Detectron2 CUDA Compiler: not available".
  - "invalid device function" or "no kernel image is available for execution".
  - Undefined CUDA symbols or cannot open libcudart.so.
  - "ImportError: cannot import name '\_C'".
  - ONNX conversion segfault after some "TraceWarning".
- See the downside of the webpage for solution.  
<https://github.com/facebookresearch/detectron2/blob/master/INSTALL.md>

## Detectron2 is ...

- A library / research platform for object detection
  - Includes everything we just learned
  - And helps us build more
- In this tutorial:
  - What's in Detectron2
  - How to Use Detectron2 (beginners)
  - How to Use Detectron2 (advanced users)

# Family of Detection/R-CNN Codebase





## Why PyTorch (Eager Mode) for detection

- Tensors of irregular shapes
- Loops / control flow / heuristics
- More hackable for researcher

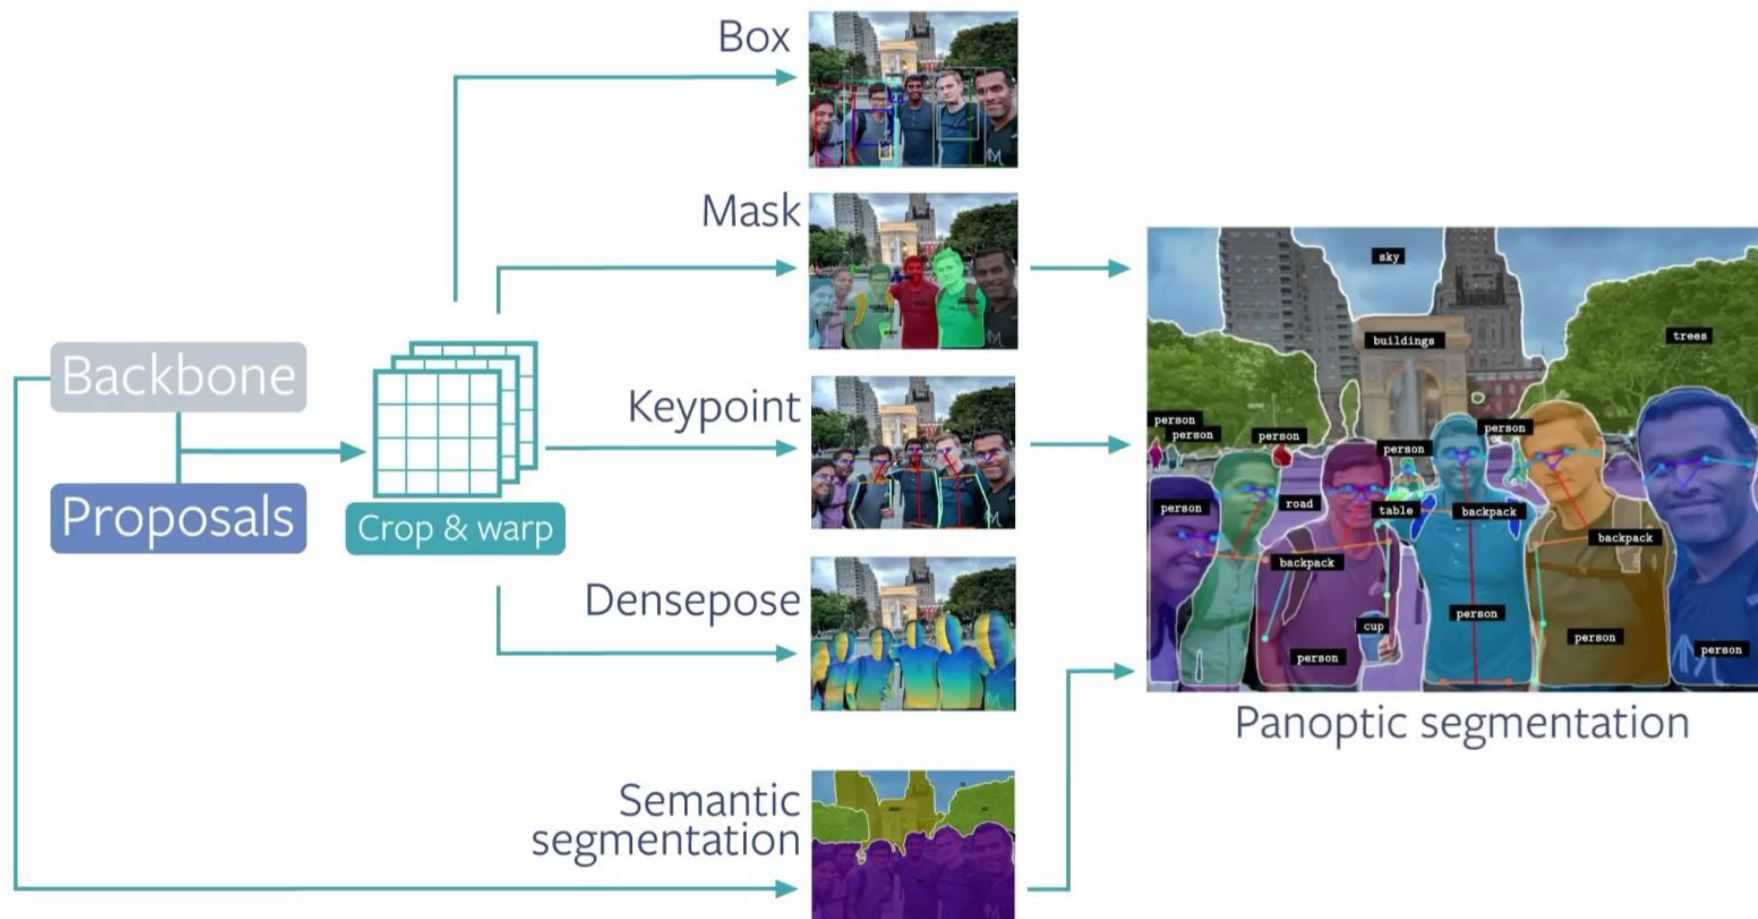


Detectron2



Detectron2

# What's in Detectron2: Generalized R-CNN Models



# What's in Detectron2

- Datasets:
  - COCO
  - LVIS
  - CityScapes
  - PascalVOC
- Tasks (data & evaluation):
  - (Rotated) Box
  - {Instance,Semantic,Panoptic} Segmentation
  - Person Keypoint, DensePose

# What's in Detectron2: MODEL ZOO

- different models for users to play with.
- baselines for researchers.

Faster R-CNN:

Name	lr sched	train time (s/iter)	inference time (s/im)	train mem (GB)	box AP	model id	download
<a href="#">R50-C4</a>	1x	0.593	0.110	4.8	35.7	137257644	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R50-DC5</a>	1x	0.380	0.068	5.0	37.3	137847829	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R50-FPN</a>	1x	0.210	0.055	3.0	37.9	137257794	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R50-C4</a>	3x	0.589	0.110	4.8	38.4	137849393	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R50-DC5</a>	3x	0.378	0.073	5.0	39.0	137849425	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R50-FPN</a>	3x	0.209	0.047	3.0	40.2	137849458	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R101-C4</a>	3x	0.656	0.149	5.9	41.1	138204752	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R101-DC5</a>	3x	0.452	0.082	6.1	40.6	138204841	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R101-FPN</a>	3x	0.286	0.063	4.1	42.0	137851257	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">X101-FPN</a>	3x	0.638	0.120	6.7	43.0	139173657	<a href="#">model</a>   <a href="#">metrics</a>

RetinaNet:

Name	lr sched	train time (s/iter)	inference time (s/im)	train mem (GB)	box AP	model id	download
<a href="#">R50</a>	1x	0.200	0.062	3.9	36.5	137593951	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R50</a>	3x	0.201	0.063	3.9	37.9	137849486	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">R101</a>	3x	0.280	0.080	5.1	39.9	138363263	<a href="#">model</a>   <a href="#">metrics</a>

RPN &amp; Fast R-CNN:

Name	lr sched	train time (s/iter)	inference time (s/im)	train mem (GB)	box AP	prop. AR	model id	download
<a href="#">RPN R50-C4</a>	1x	0.130	0.051	1.5		51.6	137258005	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">RPN R50-FPN</a>	1x	0.186	0.045	2.7		58.0	137258492	<a href="#">model</a>   <a href="#">metrics</a>
<a href="#">Fast R-CNN R50-FPN</a>	1x	0.140	0.035	2.6	37.8		137635226	<a href="#">model</a>   <a href="#">metrics</a>

COCO Instance Segmentation Baselines with Mask R-CNN

Name	lr sched	train time (s/iter)	inference time (s/im)	train mem (GB)	box AP	mask AP	model id	download
------	----------	---------------------	-----------------------	----------------	--------	---------	----------	----------

# Quick Start with Model Zoo: Inference

- Pick a model: Config + Checkpoint
- Run it:

```
cfg = get_cfg()
cfg.merge_from_file("./configs/.../mask_rcnn_R_50_FPN_3x.yaml")
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5 # set threshold for this model
# Find a model from detectron2's model zoo. You can either use the
# https://dl.fbaipublicfiles.... url, or use the following shorthand
cfg.MODEL.WEIGHTS = "detectron2://.../mask_rcnn_R_50_FPN_3x/...pkl"
predictor = DefaultPredictor(cfg)
predictions = predictor(img)
```

# Quick Start with Model Zoo: Visualization

- Show outputs on images

```
from detectron2.utils import Visualizer
vis = Visualizer(img, coco_metadata)
output = vis.draw_instance_predictions(predictions["instances"])
cv2.imshow("", output.get_image())
```

# Quick Start with Model Zoo: Visualization



# Quick Start with Model Zoo: Inference

- Show outputs on images

```
from detectron2.utils import Visualizer
vis = Visualizer(img, coco_metadata)
output = vis.draw_instance_predictions(predictions["instances"])
cv2.imshow("", output.get_image())
```

- .. And videos (with tracking)

```
from detectron2.utils import VideoVisualizer
vis = VideoVisualizer(coco_metadata)
for frame, predictions in stream_of_predictions():
    output = vis.draw_instance_predictions(frame, predictions["instances"])
    cv2.imshow("", output.get_image())
```



## Detectron2 for Advanced Users

- Write new models / Improve existing models
  - New data processing
  - Define new tasks/metrics
- 
- Maintainability
  - Deployment

## Add Something New

- Hack inside the code: `vim detectron2`
  - Quick & flexible prototyping
  - Not scalable / maintainable
- Extend existing code: `import detectron2`
  - Some\* code duplication
  - Take some time
  - Maintainable

Good research codebase should be  
Hackable and Extensible

# Hackable

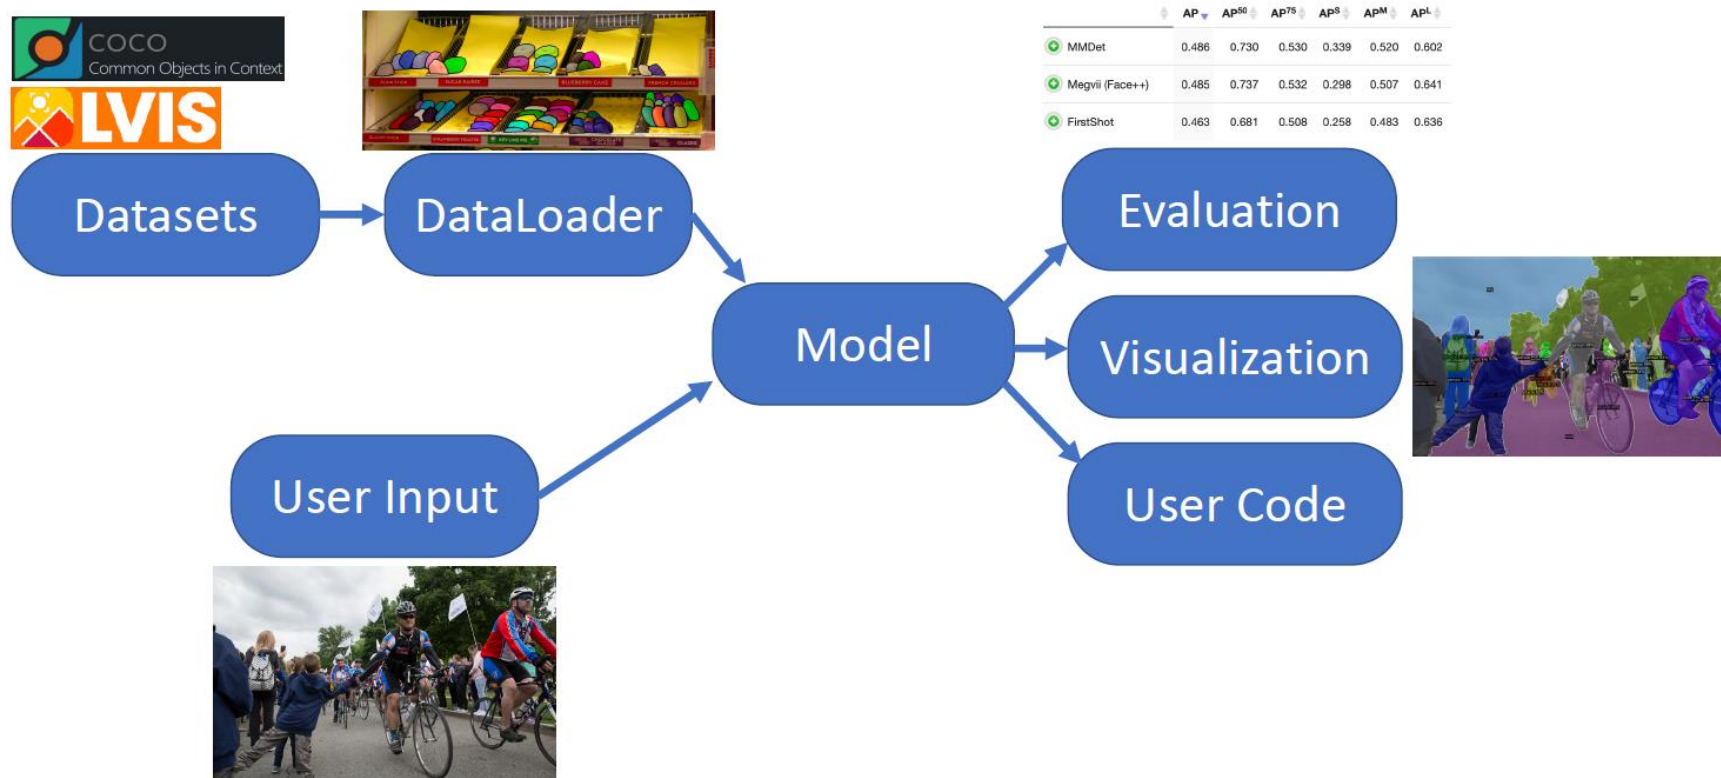
- Simple abstraction
- Straightforward implementation
- Well-documented
- ...even can enable global cfg

```
detectron2/  
├── checkpoint  
├── config  
├── data  
│   ├── datasets  
│   ├── samplers  
│   └── transforms  
├── engine  
├── evaluation  
├── layers  
├── modeling  
│   ├── backbone  
│   ├── meta_arch  
│   ├── proposal_generator  
│   └── roi_heads  
├── solver  
├── structures  
└── utils
```

## Extensible / Customizable

- Allow users to plug in custom
  - Models
  - Datasets
  - Data loading routines
  - Augmentations
  - Tasks/Metrics
  - Training logic
  - ...

# Structure of Modules



# Training

- Either write your own training loop
- OR:
- Use a built-in trainer + hooks:
  - Minimal simple training needs:
    - Model, Dataloader, Optimizer SimpleTrainer
  - Standard training also needs:
    - Evaluation / Checkpoint / LRSchedule / ... DefaultTrainer
    - Many users just want this standard workflow

`tools/plain_train_net.py` vs. `tools/train_net.py`



## Detectron2 for PA1



# Detectron2

<https://github.com/facebookresearch/detectron2>

# Training & Evaluation in Command Line

We provide a script `train\_net.py`, which is made to train configs.

- To train a model with `train\_net.py`,

```
python train_net.py --num-gpus 4 \  
  --config-file configs/R_50_1x.yaml
```

- you may need to change some parameters, e.g.:

```
python train_net.py --config-file configs/R_50_1x.yaml \  
  --num-gpus 1 SOLVER.IMS_PER_BATCH 4 SOLVER.BASE_LR 0.025
```

- To evaluate a model's performance, use

```
python train_net.py --config-file configs/R_50_1x.yaml \  
  --eval-only MODEL.WEIGHTS /path/to/checkpoint_file
```

For more options, see `./train_net.py -h`.



## D2 Hackable: Config

### detelectron2.config

Detectron2's config system uses `yaml` and `yacs`. In addition to the basic operations that access and update a config, we provide the following extra functionalities:

1. The config can have `_BASE_`: `base.yaml` field, which will **load a base config first**. Values in the base config will be overwritten in sub-configs, if there are any conflicts.
2. We provide config versioning, for backward compatibility. If your config file is versioned with a config line like `VERSION: 2`, `detectron2` will still recognize it even if we rename some keys in the future.

**Use Configs:** Some basic usage of ``CfgNode`` object is shown below:

```
from detectron2.config import get_cfg
cfg = get_cfg()      # obtain detectron2's default config
cfg.xxx = yyy        # add new configs for your own custom components
cfg.merge_from_file("my_cfg.yaml")      # load values from a file
cfg.merge_from_list(["MODEL.WEIGHTS", "weights.pth"]) # can also load values from a list of str
print(cfg.dump())    # print formatted configs
```

## D2 Hackable: Config

### `detecelectron2.config`

- The default configurations for D2 that loaded at the very first:
  - `detecelectron2/config/defaults.py`
- For PA1, the default configurations for FCOS:
  - `fcos/config/defaults.py`
- You can control such hyperparameters on:
  - ``fcos/config/defaults.py``, or
  - ``configs/*.yaml``

## D2 Hackable: Datasets

### detectron2.data

- Datasets are assumed to exist in a directory './datasets' relative to your current working directory.
- Under this directory, detectron2 expects to find datasets in the following structure:

```
coco/  
  annotations/  
    instances_{train,val}2017.json  
    person_keypoints_{train,val}2017.json  
  {train,val}2017/  
    # image files that are mentioned in the corresponding json
```

## D2 Hackable: Datasets

detecetron2.data

### Standard Dataset Dicts

- For standard tasks (detection, instance segmentation), we load the original dataset into *list[dict]* with a specification similar to COCO's json annotations. This is our standard representation for a dataset.
- Each dict contains information about one image. The dict may have the following fields.

*file\_name*: the full path to the image file.

*height*, *width*: integer. The shape of image.

*image\_id* (str or int): a unique id that identifies this image. Used during evaluation to identify the images

*annotations* (list[dict]): each dict corresponds to annotations of one instance in this image. Each dict may contain the following keys:

*bbox* (list[float]): list of 4 numbers representing the bounding box of the instance.

*bbox\_mode* (int): the format of bbox. It must be a member of `structures.BoxMode`.

Currently supports: `BoxMode.XYXY_ABS`, `BoxMode.XYWH_ABS`.

*category\_id* (int): an integer in the range  $[0, \text{num\_categories})$  representing the category label. The value *num\_categories* is reserved to represent the *"background" category*.

## D2 Hackable: Models

### detectron2.modeling

- Models (and their sub-models) in detectron2 are built by functions such as `build_model`, `build_backbone`, `build_roi_heads`.
- Note that `build_model` only builds the model structure, and fill it with random parameters. To load an existing checkpoint to the model, use `DetectionCheckpointer(model).load(file_path)`.
- Detectron2 recognizes models in pytorch's `.pth` format, as well as the `.pkl` files in model zoo.
- When loading pre-trained FCOS checkpoint file, use `fcos.checkpoint.AdetCheckpointer`.
- You can use a model by just `outputs = model(inputs)`.

# Model Input Format

- All builtin models take a `list[dict]` as the inputs. Each dict corresponds to information about one image.
- The dict may contain the following keys:
  - “image”: `Tensor` in (C, H, W) format. The meaning of channels are defined by `cfg.INPUT.FORMAT`. Image normalization, if any, will be performed inside the model.
  - “instances”: an `Instances` object, with the following fields:
    - “gt\_boxes”: a `Boxes` object storing N boxes, one for each instance.
    - “gt\_classes”: `Tensor` of long type, a vector of N labels, in range `[0, num_categories)`.
    - “gt\_masks”: a `PolygonMasks` or `BitMasks` object storing N masks, one for each instance.
  - “height”, “width”: the **desired** output height and width, which is not necessarily the same as the height or width of the `image` input field. For example, the `image` input field might be a resized image, but you may want the outputs to be in **original** resolution.

## Model Output Format

- When in **training mode**, the builtin models output a `dict[str->ScalarTensor]` with including all the losses.
- When in **inference mode**, the builtin models output a `list[dict]`, one dict for each image. Each dict may contain the following fields:
  - “instances”: `Instances` object with the following fields:
    - “pred\_boxes”: `Boxes` object storing N boxes, one for each detected instance.
    - “scores”: `Tensor`, a vector of N scores.
    - “pred\_classes”: `Tensor`, a vector of N labels in range `[0, num_categories)`.
    - “pred\_masks”: a `Tensor` of shape `(N, H, W)`, masks for each detected instance.

## D2 Hackable: Structures → Instances

### `detectron2.structures.Instances`

- Represents a list of instances in an image. It stores the attributes of instances (e.g., boxes, masks, labels, scores) as “fields”.
- Set/Get a field:

```
instances.gt_boxes = Boxes(...)  
print(instances.pred_masks) # a tensor of shape (N, H, W)  
print('gt_masks' in instances)
```

- **`len(instances)`**: returns the number of instances.
- Indexing: **`instances[indices]`** will apply the indexing on all the fields and returns a new **Instances**. Typically, **`indices`** is a integer vector of indices, or a binary mask of length **`num_instances`**,



## D2 Hackable: Structures → Boxes

### `detectron2.structures.Boxes`

- Stores a list of boxes as a Nx4 `torch.Tensor`. It supports some common methods about boxes (*area*, *clip*, *nonempty*, etc), and also behaves like a Tensor (support indexing, *to(device)*, *.device*, and iteration over all boxes)
- Set a field: `tensor (Tensor[float])` – a Nx4 matrix. Each row is (x1, y1, x2, y2).

```
__init__(tensor: torch.Tensor)
```

## D2 Hackable: Structures → ImageList

### `detectron2.structures.ImageList`

- Holds a list of images (of possibly varying sizes) as a single tensor. This works by padding the images to the same size, and storing in a field the original sizes of each image.

`image_sizes`: each tuple is (h, w)

- Type: `list[tuple[int, int]]`

`__init__(tensor: torch.Tensor, image_sizes: List[Tuple[int, int]])`

- **Parameters:**

**tensor** (Tensor) – of shape (N, H, W) or (N, C<sub>1</sub>, ..., C<sub>K</sub>, H, W) where K ≥ 1

**image\_sizes** (list[tuple[int, int]]) – Each tuple is (h, w). It can be smaller than (H, W) due to padding.

`__getitem__(idx: Union[int, slice]) → torch.Tensor`

- Access the individual image in its original size.
- Returns: *Tensor*- an image of shape (H, W) or (C<sub>1</sub>, ..., C<sub>K</sub>, H, W) where K ≥ 1



# FCOS Implementation

# Programming Assignment 1

- We provide detectron2-based skeleton code for FCOS detector.
- The followings are prepared in the skeleton code:
  - Configurations
    - skeleton/configs/\*.yaml*
    - skeleton/fcos/config/defaults.py*
  - Meta-architecture containing backbone structure for FCOS
    - skeleton/fcos/modeling/{backbone, meta\_arch}/*
  - Any other utilities related to operating FCOS detector
    - Skeleton/train\_net.py*
    - Skeleton/fcos/{checkpoint, evaluation, layers, utils}/*
  - **Skeleton code** for FCOS pipelines, including FCOS heads
    - skeleton/fcos/modeling/fcos/{fcos, fcos\_head, fcos\_losses, fcos\_targets}.py*
- All you have to do is to complete each part of skeleton code step-by-step.

# Programming Assignment 1

- Followings are **requirements** for you to earn all credits (100) for PA1.
  - Implementation of baseline FCOS detector.
  - Train the baseline model on COCO dataset. (options for ``configs/R_50_1x.yaml``)
  - Evaluate on COCO validation set, and record Average Precisions on your report.
  - Report qualitative detection results on some selected COCO validation images, or on any other images, using `:class:`detectron2.utils.Visualizer``.
- **Do not change any other options in config files.**
  - Especially, we fix backbones and input resolutions for reproduction purpose.
  - You may change configs for ``SOLVER`` when following linear scaling rule [1].

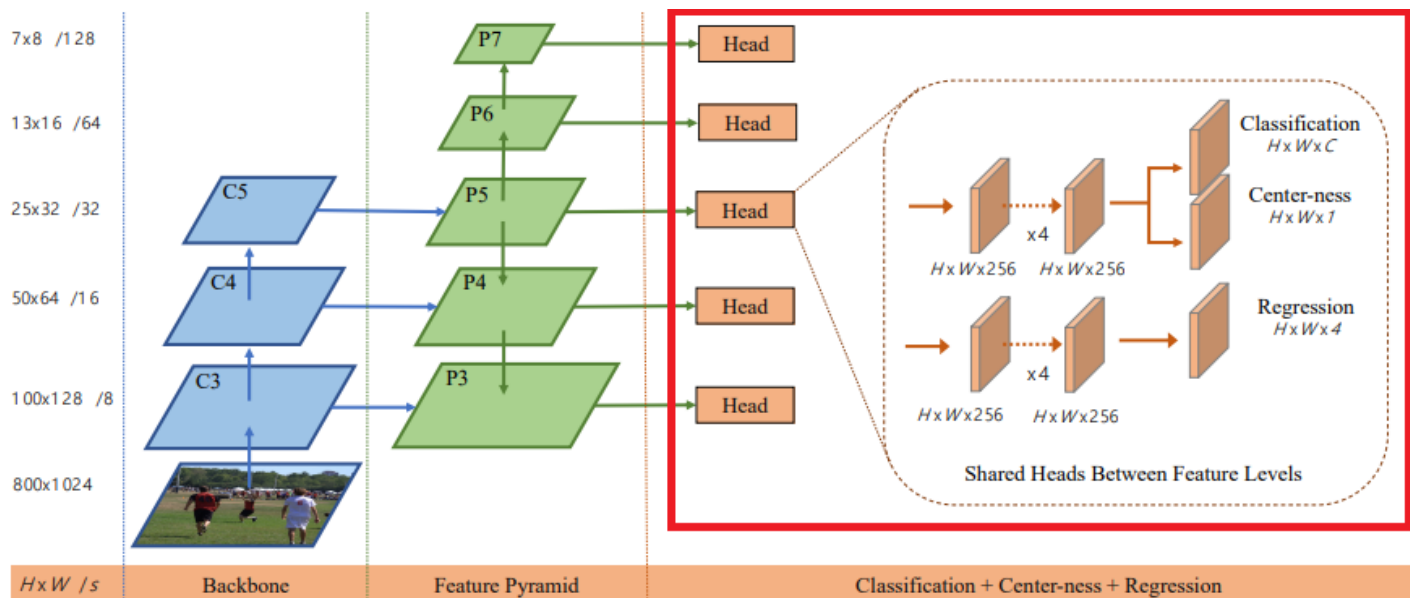
## Training FCOS

- Default hyperparameters related to training schedules are set assuming that we train the model with 4-GPUs, with total batch size 16. (4 images per GPU).
- If you need to change the total batch size by some reasons (i.e., #GPUs, #imgs/GPU), you should **modify** the initial learning rate value and total training iterations (including what iterations to down-scale the learning rate), following the **linear scaling rule** specified in [1].
- `SOLVER` hyperparameters are specified in `configs/Base-FCOS.yaml`.

# Step 1. FCOS Head

`fcos.modeling.fcos.FCOSHead`

- FCOS head structure is as following.
  - The heads across all different feature levels share the same parameters.
- Please refer to **Figure 2** and **"Network Outputs"** on Section 3 from the FCOS paper.
- To make clear, for the final results from each of classification, centerness, and regression layer, **applying activation function (i.e., sigmoid) is not needed**. Proper activation function is applied when computing loss function and decoding bounding box proposals.



## Step 1. FCOS Head

### `fcos.modeling.fcos.FCOSHead`

- Tips

- `:meth:~FCOSHead.__init__` is implemented.
- You need to implement the remaining methods:  
`:meth:~FCOSHead.{_init_layers, _init_weights, forward}`.
- When you define a conv2d layer, you may use `detectron2.layers.Conv2d` wrapper.
- For weight initialization in `:meth:~FCOSHead._init_weights`, you may use `:func:~fcos.layers.normal_init`.
- Make sure that you control each building components by the hyperparameters specified in config file automatically.



## Step 2. FCOS Targets

### `fcos.modeling.fcos.FCOSTargets`

- Here, you need to compute ground-truth targets needed for training FCOS:
  - Finally returns ``labels`` and ``bbox_targets``.
- The codeflow is constructed as following.
  - `:func:`get_points``: compute (x, y) coordinates for feature points across all feature levels back-projected into image coordinates.
  - `:func:`fcos_target``: compute *`class labels`* and *`regression targets`* for every feature points computed in above, across all feature levels.
  - `:func:`compute_centrerness_targets``: compute centerness targets for every feature points, given bbox targets. This function is called from `:func:`FCOSLosses``.

## Step 2-1. FCOS Targets: `get_points`

### `FCOSTargets.get_points`

- `:func:`get_points`` evaluates `:func:`get_points_single`` and accumulates the results per feature level.
- Complete the code of `:func:`get_points_single``.
- Hint: using `:func:`torch.meshgrid`` will be helpful.
- From FCOS paper **Section 3.1**,  
For each location  $(x, y)$  on feature map  $F_i$ ,  
we can map back onto the input image as  $\left(\left\lfloor \frac{s}{2} \right\rfloor + xs, \left\lfloor \frac{s}{2} \right\rfloor + ys\right)$

## Step 2-2. FCOS Targets: `fcos_target`

### `FCOSTargets.fcos_target`

- `:func:`fcos_target`` evaluates `:func:`fcos_target_single_image`` and accumulates the results per image.
- Implement `:func:`fcos_target_single_image``.
- Since the process of obtaining classification targets and regression targets is not intuitive, we provide line-by-line skeleton code examples.
- Replace the *'NotImplemented'* with your answers.
- Or, you may re-write the whole codes from scratch on yourself.
- We recommend you to read **Section 3.2** on FCOS paper.

## Step 2-3. FCOS Targets: `compute_centerness_targets`

`FCOSTargets.compute_centerness_targets`

- Given regression targets, a tensor shape of  $(N, 4)$ , compute the centerness targets.
- See **Equation (3)** and **Figure 3** from FCOS paper.

## Step 3. FCOS Losses

`fcos.modeling.fcos.FCOSLosses`

- FCOS adopts two types of loss functions: ``focal loss`` and ``IOU loss`` for classification and regression task, respectively.
- For detail, see **Section 3.1 “Loss Function”** on FCOS paper.
- We provide loss function implementations, so you are free from implementing loss functions.

## Step 4. Inference

### FCOS.predict\_proposals

- In inference phase, :meth:`FCOS.predict\_proposals` predicts a complete form of detected bounding boxes with corresponding class indices from the outputs of FCOS head.
- This process includes:
  - **Ranking proposals** (top-k) and sorting based on class scores.
  - **Score thresholding** to filter out low-quality predictions.
  - **Decoding regressions** (l, t, r, b) into bounding boxes with corresponding classes.
  - **Non-maximum suppression** for removing duplicated boxes.
- See **Section 3.3** on FCOS paper describing the inference phase.

## Step 4. Inference

### FCOS.predict\_proposals

- Here are some tips for implementing inference phase for FCOS detector.
  - :meth:`FCOS.predict\_proposals` evaluates :meth:`FCOS.predict\_proposals\_single\_image` and accumulates the results per image. You need to complete this per-image operation.
  - When ranking the proposals, class scores should be multiplied by centerness scores.
  - One trick for ranking proposals is that for a single row of `scores` variable from the code, it can be regarded as **total `K` number of independent predictions for all `K` classes with their scores**, rather than only `one prediction` per row with the maximum score as the most likely class as the same with standard R-CNN family detectors, since the **activation of classification logits is sigmoid function**.
  - Inside the per feature-level loop, the predicted bounding boxes should be decoded in the form of (x1, y1, x2, y2), a tensor shape of (N, 4).
  - For the very final classification scores, applying **square root** is needed to prevent down-weighted scores (caused by multiplication with centerness scores).
  - You should utilize :class:`Instances` and :class:`Boxes` to properly construct per feature-level model outputs with the right format.
  - This processed :class:`Instances` object that contains detection results are accumulated across all feature levels and then forwarded to nms-layer which is given.

## Extra Credits

- You can earn extra credits if you implement methods corresponding to **“Improvements” on Table 3** on FCOS paper, **including ‘w/ GN’ option** and make comparisons quantitatively/qualitatively with the baseline model. (For GN, use `:class: `torch.nn.GroupNorm``.)
- When training the improved version, use ``configs/R_50_1x_improve.yaml``.
  - You may freely modify the options for the improvements as far as you implemented.
- You may compare the only your final version of the improved model with the baseline if your GPU environment is limited.



## ExtraExtra Credits

- Expand the detection task to Instance Segmentation task.
- FCOS detector as a proposal generator for `GeneralizedRCNN` framework.
- Attach a ROI pooling layer (i.e., ROIAlign) and a standard mask head implemented in Detectron2 in a cascading manner after FCOS bbox detections.
  - Use FCOS detection results for sampling ROIs that will be provided to `mask roi head`.
  - Also, regard those detection results as the final bbox results. (No additional bbox head)
- In this part, you should generate proper configuration file for training. Also making some modifications on original D2 codes will be needed.
  - For modification of codes, you write a new custom model code that overrides the existing relevant components (i.e., ROIHeads, MaskHead, Sampler, ...) from D2, and make changes on your code wherein only required.
- When FCOS-based instance segmentation task is implemented,
  - Report Mask Average Precisions
  - Compare with standard `R-50-FPN Mask R-CNN` model quantitatively/qualitatively.
  - You may download the `COCO-pretrained Mask R-CNN weights` from D2 model zoo and make comparisons with your implemented Instance Segmentation model.

## Cautions and Comment

- **Do not modify any files or folders** except for python files located in the path: ``skeleton/fcos/modeling/fcos/*.py``
- Exceptionally, you may define helper functions on existing files or new ``*.py`` file for your convenience.
- There will be partial points on every steps, so try to do your best.
- Collaboration with other students is prohibited. Please do it yourself and comment on your code in order to show your understanding.

## PA1 Submission

- **Due: May 12<sup>th</sup> 11:59PM** (No deadline extension)
- **To:** Youngtaek Oh (youngtaek.oh [at] kaist.ac.kr)
- **Submission should include:**
  - Report (`.docx` or `.pdf`)
  - Source code (do not include dataset folder)
  - Only one `.pth` checkpoint that yields the best AP.  
You should mention which method your result corresponds to in Table 3.
  - Zip your all inclusions named as `PA1\_{StudentID}\_{NAME}.zip`,  
i.e., `PA1\_20193358\_YoungtaekOh.zip`.
- **The report should include:**
  - Your understanding of each step of FCOS implementation.
  - Reproduced AP for the baseline model (on Table 3, 4<sup>th</sup> row of FCOS paper)
  - Visualization of detection results from the baseline model.
  - (Extra) Average Precisions for your improved model (Table 3 Improvements) and comparisons with the baseline model.
  - (ExtraExtra) The mask Average Precisions of your Instance Segmentation model and comparisons with `R-50-FPN Mask R-CNN` model.

## Revision History

Revision	Date	Description
0	2020-04-14	Release
1	2020-04-26	p. 39, clarification of FCOS head predictions.
2	2020-05-01	1. Fixing a bug on :func: <code>`fcos_target_single_image`</code> in <code>`fcos_targets.py`</code> 2. Fixing typo on docstring.